

Using an LLM to improve prompts to an LLM

References

- [LLMs are Human-Level Prompt Engineers \(https://arxiv.org/pdf/2211.01910.pdf\)](https://arxiv.org/pdf/2211.01910.pdf)
 - [APE summary \(https://sites.google.com/view/automatic-prompt-engineer\)](https://sites.google.com/view/automatic-prompt-engineer)

The Automatic Prompt Engineer (APE) is a system to *improve* upon prompts

- given a prompt
- APE will create a prompt that is *more effective*

It uses an LLM for multiple purposes

- to create variations of the given prompt
- to evaluate which variation is best

Using APE to improve upon Instruction Following

APE has been demonstrated to improve prompts **for a specific task** (not improvement of general prompts)

The task is: creating *instructions*

- to use in fine-tuning a raw LLM into an helpful Assistant
- we described Instruction Following in the module [Synthetic data for Instruction Following \(LLM Instruction Following Synthetic Data.ipynb\)](#).

In order to create an example of Instruction Following behavior we need a prompt with multiple parts

- a textual *task description (instruction)* that describes a task to be performed
- zero or more exemplars: demonstrating the input/output relationship described by the instruction

Given just the exemplars (the second part)

- we want APE to *create* the "best" instruction (first part)

The APE method for automatically generating "good" instructions is conceptually simple.

As a first step, we get an LLM to generate plausible instructions.

- Given the exemplars
- Create a prompt
- Whose "response" is an *instruction* that is a plausible description of the input/output relation in the exemplars

We do this several times to generate multiple plausible instructions, conditional on the exemplars.

In a second step, we rank the multiple instruction candidates created in the first step.


- we create a prompt requesting that the LLM *rank* the instructions created in the first step
- we filter the instructions to the most highly ranked results

As an optional third step, we can improve upon the *diversity* of the instructions selected in Step 2

- create a prompt requesting that the LLM generate a *variation* of a selected instruction
- Use an LLM to create a distribution of instruction, conditional on the exemplars

Here is a picture of the workflow.

APE Workflow



(a) Automatic Prompt Engineer (APE) workflow

Attribution: <https://arxiv.org/pdf/2211.01910.pdf#page=2>

Step 1: LLM as Inference Model

Goal: Get the LLM to create multiple *instruction* candidates, given the intended response

- intended *response* are the exemplars demonstrating the input/output relation
 - Input: prove, Output: disprove

Given the context (response), the LLM is prompted to generate an *instruction* that could cause the given response

- prompt is a Masked Language Modeling task: fill in the mask (<INSERT>)

This is in the same spirit as [Backtranslation](#)

([LLM Instruction Following Synthetic Data.ipynb#Instruction-Backtranslation](#)).

- learns a new model to map from response to instruction
- rather than adapting an LLM to do the same

Steps 2 and 3: LLM as Scoring Model

Step 1 has created multiple possible instructions that are consistent with the exemplars.

We wish to rank them.

The ranking is performed by prompting the LLM to compute the likelihood

- that a given response (an exemplar)

Input: direct, Output: indirect

- is consistent with each candidate instruction

The candidate instructions are ranked from highest Likelihood to lowest.

Note

Likelihood is expressed as the log probability

- is a negative number since probabilities are fractions
 - less negative numbers are higher probabilities
- is the probability of the generated sequence
 - product of the individual probabilities of the tokens in the sequence

Steps 4 (Optional): LLM as Re-sampling model

Here, we create multiple variants of a highly ranked instruction.

Given the candidate Instruction selected by the previous step

- generate a variation of the instruction
- by asking the LLM to create it via text completion

Generate a variation of the following instruction ...

Input: write the antonym of the word; Output: <COMPLETE>

Each step is implemented as an instance of the pre-trained LLM's ability to complete text (or fill in a mask).

No fine-tuning or adaptation of weights is involved.

Forward/Reverse mode generation of candidates

This is a minor technical point.

The prompt in Step 1 of the workflow above is not in the format consistent with text-continuation

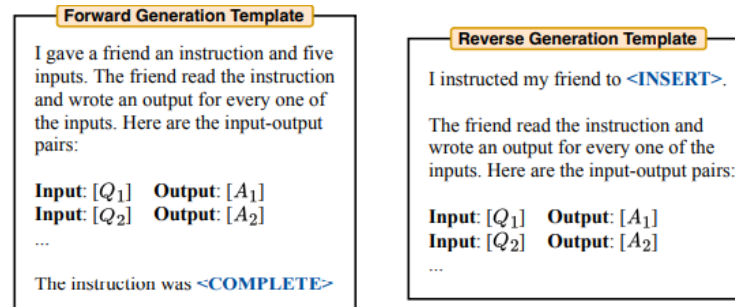
- format is called *forward generation*
- so must use an LLM that solves Masked Language task, rather than text-continuation

An alternate prompt can be used that is consistent with text-continuation

- format is called *reverse generation*

APE Forward/Reverse Generation templates

Attribution: <https://arxiv.org/pdf/2211.01910.pdf#page=4>



APE evaluation: super-human performance !

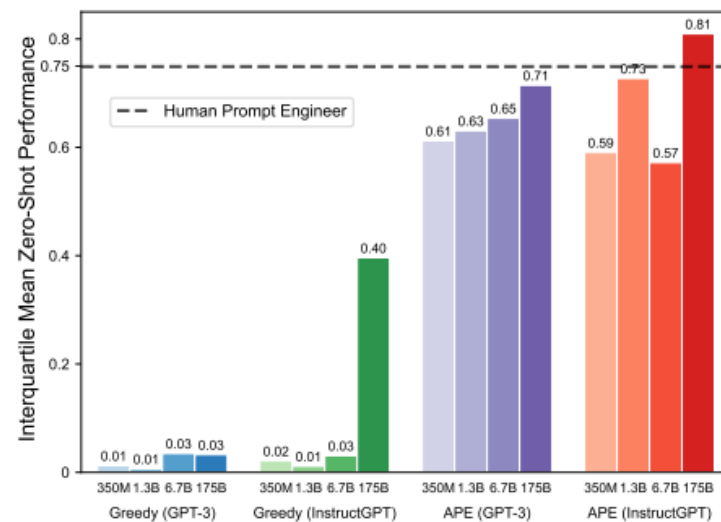
Here is a comparison of APE generated prompts

- versus
 - an alternate method (previously published), labeled "Greedy"
 - a human engineer (horizontal dotted line)
- evaluated on models of various sizes
 - GPT-3
 - Instruct GPT-3 (fine-tuned for instruction following)
- using 24 NLP tasks

Note

The reported statistic is *interquartile mean* (i.e., average after dropping the upper and lower 25% of results)

APE Workflow and Results



(b) Interquartile mean across 24 tasks

Attribution: <https://arxiv.org/pdf/2211.01910.pdf#page=2>

Zero-shot: Improving on "Let's think step by step"

[Chain of Thought \(CoT\) prompting \(NLP Beyond LLM.ipynb#Chain-of-thought-prompting\)](#).

- is a simple technique
- for create prompts with better performance
- for multi-step reasoning problems

In the few-shot setting

- exemplars demonstrate step by step reasoning
- eliciting the LLM to produce text continuation that also exhibits step by step reasoning

In the zero-shot setting, it simply involves appending

Let's think step by step

to the prompt

Chain of Thought Prompting




Figure 1: Example inputs and outputs of GPT-3 with (a) standard Few-shot ([Brown et al., 2020]), (b) Few-shot-CoT ([Wei et al., 2022]), (c) standard Zero-shot, and (d) ours (Zero-shot-CoT). Similar to Few-shot-CoT, Zero-shot-CoT facilitates multi-step reasoning (blue text) and reach correct answer where standard prompting fails. Unlike Few-shot-CoT using step-by-step reasoning examples **per task**, ours does not need any examples and just uses the same prompt “Let’s think step by step” *across all tasks* (arithmetic, symbolic, commonsense, and other logical reasoning tasks).

Attribution: <https://arxiv.org/pdf/2201.11903.pdf>

Let's use APE (<https://arxiv.org/pdf/2211.01910.pdf#page=19>) to find a zero-shot prompt appendage that improves upon

Let's think step by step

That is: the instruction consists of

- the task description
- followed by a "magic suffix" (e.g., "let's think step by step")

We want to find the best "magic suffix".

The authors use the following template (where INPUT and OUTPUT are place-holders for an actual question and answer pair).

Instruction: Answer the following question

Q: INPUT

A: Let's <INSERT> OUTPUT

We are using forward-mode generation to get APE

- to create a phrase that follows the INPUT
- that begins with the word "Let's"

APE creates

Let's work this out in a step by step way to be sure we have the right answer.

and the author's demonstrate improved performance on several benchmarks.

This is a nice demonstration of using an LLM to help craft better prompts to LLM's.

In [2]: `print("Done")`

Done

