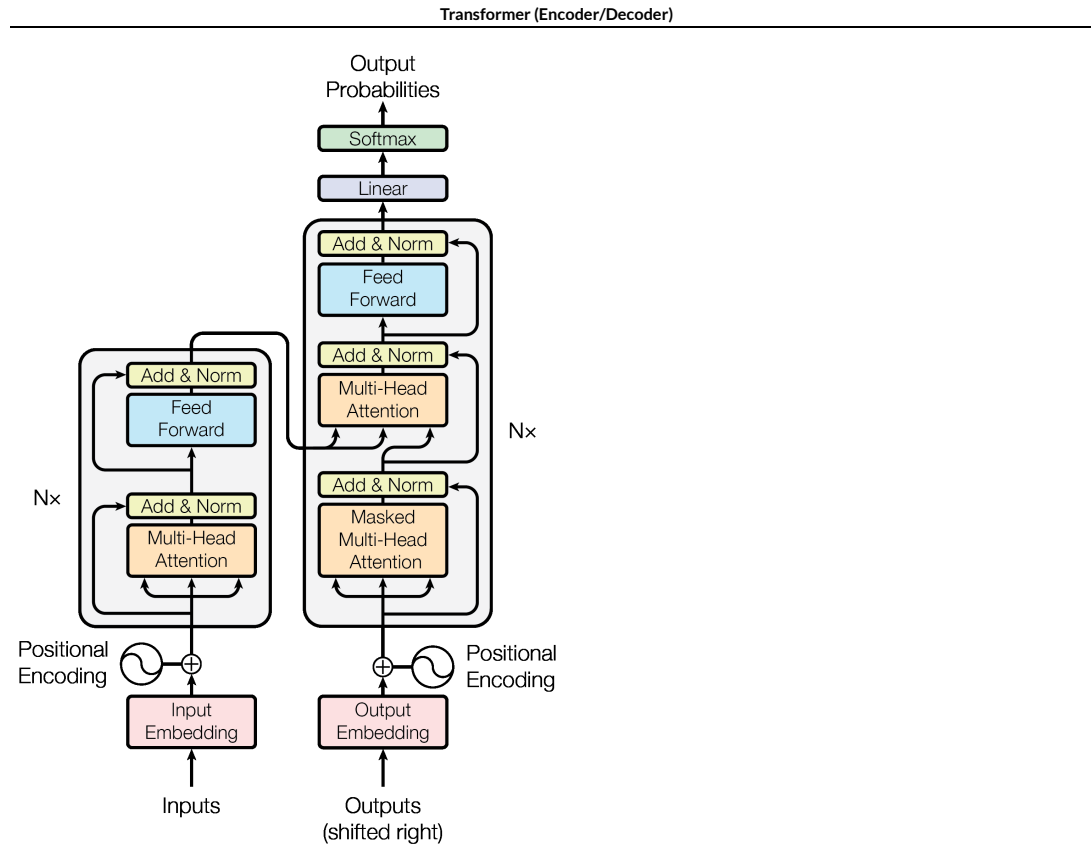


# Transformer: Intuition

We try to briefly explain what each the "moving parts" of the Transformer is doing.





# General

Each of the paths in the Transformer is a vector of length  $d_{\text{model}}$

- sometimes just referred to as  $d$

Having a common length simplifies the architecture

- can stack Transformer blocks (since input and output are same size)
- Self-Attention and Cross-Attention:
  - map a query of size  $d$  to an output of size  $d$
- Needed for the Residual Connection (Add and Norm)
  - adding the input of Attention to the output of Attention
    - need to be same length

## Residual connections

- [Residual connections from Intro course \(RNN Residual Networks.ipynb\)](#).

# Encoder

The Encoder style Transformer can either be used

- stand-alone: create a fixed length, alternate length representation of the input
  - for further processing: e.g., Classification
- as part of an Encoder-Decoder architecture
  - transform the Input sequence  $\mathbf{x}$  in a *processed sequence*  $\bar{\mathbf{h}}_{(1)}, \dots, \bar{\mathbf{h}}_{(\bar{T})}$
  - to be consumed by a Decoder


## Encoder Self-Attention

The Self-Attention allows each  $\bar{\mathbf{h}}_{(\bar{t})}$  to depend on the **complete** input sequence  $\mathbf{x}$ .

If we view  $\bar{\mathbf{h}}_{(\bar{t})}$  as the "meaning" of  $\mathbf{x}_{(t)}$

- it is a meaning based on the *full context*
  - not just the preceding elements  $\mathbf{x}_{(1:t-1)}$

By making the meaning dependent on the full context, we can disambiguate the meaning of the word "it"



## Decoder

The salient characteristic of a Decoder style Transformer is the *autoregressive* behavior

- generates an output one token at a time
- by appending each generated token to the sequence of already-generated tokens

The Decoder style Transformer can either be used

- stand-alone: for generative tasks
- as part of an Encoder-Decoder architecture



## Decoder Self-Attention

The input needs to decide why of the previously-generated Decoder outputs (now Decoder inputs) to attend to.

There are multiple potential uses for this

- to help generate the "next" token (goal of the Decoder), by referencing the partially complete Decoder output
- to help in the Cross Attention step
  - decide which part of the Inputs  $\mathbf{x}$  (Decoder Outputs) to attend to
  - "looking up" facts, e.g., our Question Answering example or Language Translation example

Note the use of Causal Masking

- we can only reference the Decoder output already generated

## Decoder-Encoder Cross-Attention

The output of the Decoder Self-Attention is used as a "query"

- to reference the relevant part of the Input  $\mathbf{x}$

## Feed Forward Network (FFN)

Maps the output of the Decoder-Encoder Attention into the "next output token".

- actually: it is still an embedding of the next token, rather than the true next token
  - that way: it can be appended to the already-generated output to become the Decoder input for next position

This acts as a Classifier

- mapping the input
- to a vector of logits
  - one element per possible element of the Output Vocabulary

There is some evidence that

- the parameters of the FFN are where "world knowledge" is stored
  - every "fact" learned during training

# Linear

This layer is append *only* to the final block in the stacked Transformer blocks.

It acts as a typical Classifier

- "classifies" the final block's output of length  $d$
- returning a vector
  - whose length is equal to number of elements of the Vocabulary
  - each element is a logit
    - to be converted into probability distribution over elements of the Vocabulary

# Softmax

Converts the logit for each possible element of the Vocabulary

- into Probability that the element is the next Decoder Output

In [2]: `print("Done")`

Done

